

CS1101: Foundations of Programming

File handling



**Dept. of Computer Science & Engineering
Indian Institute of Technology Patna**

What is a file?

- A named collection of data, stored in secondary storage usually
- Typical operations on files:
 - Open
 - Read
 - Write
 - Close
- How is a file stored?
 - Stored as sequence of bytes, logically contiguous, need not be physically contiguous on disk
 - C gives us a simplified view of a file stored on the disk

File types

- Primarily two kinds of files
 - Text — Human readable files
 - A text editor can show the contents of a text file
 - Binary — Primarily not meant for human reading
 - Executable, video, audio, image files
 - You need special programs to view / process the file meaningfully (like, need some image viewer tool to see images, similarly video player, etc.)

File handling in C

- In C, we use FILE* to represent a pointer to file
- fopen() is used to open a file. It returns the special value NULL to indicate that it is unable to open the file

```
FILE *fptr;  
char filename[] = "myfile.dat";  
fptr = fopen(filename, "w");  
if(fptr == NULL){  
    printf("Error in file creation\n");  
    /* do something as per need */  
}
```

Modes for opening files

- The second argument in `fopen` is the **mode** in which we open the file. There are three basic modes
 - “r” opens a file for reading
 - “r+” allows write
 - “w” creates a file for writing and writes over all previous contents, deletes any previous file of the same name (**CAREFUL**)
 - “w+” allows read
 - “a” opens file for appending – writing at the end of file, previous content remain intact
 - “a+” allows read

The exit() function

- Sometime error checking means we want an immediate exit from a program
- In main() function we can use return to stop the execution of the program
- In any function we can use exit() to do this
- exit() is declared in stdlib.h
 - exit(1); // to indicate status of exit
 - exit(2);
 - exit(3);
 - ...

Usage of exit() function

```
FILE *fptr;  
char filename[]="myfile.dat";  
fptr = fopen(filename, "w");  
if(fptr == NULL){  
    printf("Error in file creation\n");  
    exit(0);  
}
```

Writing to a file using fprintf()

- fprintf() works just like the printf() except that its first argument is a file pointer

```
FILE *fptr;  
int a=3, b=8;  
char filename []="myfile.dat";  
fptr = fopen(filename, "w");  
fprintf(fptr, "First file\n");  
fprintf(fptr, "%d %d\n", a, b);
```

Reading data from a file using fscanf()

- The C library maintains a file-pointer to remember the position up to which a file has been read so far. The file-pointer moves forward with each read operation
- Next read operation (e.g., fscanf , fgets , fgetc) will give the contents of the file after this position.
- Each function for reading from a file has a way to inform that the end of file has been reached (usually by returning a special value like NULL or EOF)

```
FILE *fptr;  
int a, b;  
fptr = fopen("myfile.dat", "r");  
fscanf(fptr, "%d%d", &a, &b);
```

EOF (End of file)

- EOF is a special value that signifies that the file pointer has reached the end of the file stream.
- EOF is returned by `fgetc()` and `fscanf()` if the end of file has been reached.
- Another way to detect the end of a file is to use the call `feof(fp)`.
- `feof()` returns true only after a read operation from the file fails.
- `feof()` cannot probe and notify that the next read operation will fail. This is oftentimes not possible because whether the end-of-file is reached or not depends on what you plan to read next (like an `int` or a `char`).

Example of EOF (End of file)

- Let us assume the task is to read a file containing a set of integers and print the average of those numbers
- Consider that the last number is 10 in file input.txt and then it has few more new lines and spaces
- So, %d reading will fail but a %c will not
- You need to attempt a reading and then check for EOF or feof()

input.txt:

43 24 9 0

23 45 77

34 12

10

Checking of end-of-file using EOF

```
#include <stdio.h>
#include <stdlib.h>
int main() {
    FILE *fp; int n, sum, x;
    fp = (FILE *)fopen("input.txt", "r");
    if(fp == NULL) exit(1);
    n = sum = 0;
    while (1) {
        if (fscanf(fp, "%d", &x) == EOF) break;
        ++n; sum += x;
        printf("%d-th integer: %d\n", n, x);
    }
    fclose(fp);
    printf("Average: %.2f\n", ((float)sum)/n);
    exit(0);
}
```

Output:

1-th integer: 43
2-th integer: 24
3-th integer: 9
4-th integer: 0
5-th integer: 23
6-th integer: 45
7-th integer: 77
8-th integer: 34
9-th integer: 12
10-th integer: 10
Average: 27.70

Checking of end-of-file using EOF

```
#include <stdio.h>
#include <stdlib.h>
int main() {
    FILE *fp; int n, sum, x;
    fp = (FILE *)fopen("input.txt", "r");
    if(fp == NULL) exit(1);
    n = sum = 0;
    while (1) {
        fscanf(fp, "%d", &x); if(feof(fp)) break;
        ++n; sum += x;
        printf("%d-th integer: %d\n", n, x);
    }
    fclose(fp);
    printf("Average: %.2f\n", ((float)sum)/n);
    exit(0);
}
```

Output:

1-th integer: 43
2-th integer: 24
3-th integer: 9
4-th integer: 0
5-th integer: 23
6-th integer: 45
7-th integer: 77
8-th integer: 34
9-th integer: 12
10-th integer: 10
Average: 27.70

Checking of end-of-file using EOF: Wrong program

```
#include <stdio.h>
#include <stdlib.h>
int main() {
    FILE *fp; int n, sum, x;
    fp = (FILE *)fopen("input.txt", "r");
    if(fp == NULL) exit(1);
    n = sum = 0;
    while (!feof(fp)) {
        fscanf(fp, "%d", &x);
        ++n; sum += x;
        printf("%d-th integer: %d\n", n, x);
    }
    fclose(fp);
    printf("Average: %.2f\n", ((float)sum)/n);
    exit(0);
}
```

Output:

1-th integer: 43
2-th integer: 24
3-th integer: 9
4-th integer: 0
5-th integer: 23
6-th integer: 45
7-th integer: 77
8-th integer: 34
9-th integer: 12
10-th integer: 10
11-th integer: 10
Average: 26.09

Reading data from file using fgets()

- We can read a string from a file using fgets()
- fgets() takes 3 arguments – a string, maximum number of characters to store in the string, and a file pointer. It returns NULL if there is an error (or end of file is reached)

```
FILE *fp;
char str[1000];
fp = (FILE *)fopen("input.txt", "r");
...
while (fgets(str, 1000, fp) != NULL) {
    /* read 999 chars at most at a time */
    printf("String read is: %s\n",str);
}
...
```

Reading data from file using fgets()

- A maximum of size 1 byte will be read from the input file stream
- The reading includes the new line character if it appears in these many bytes.
- fgets() null-terminates the string by putting the null character '\0'.
- With appropriate size, fgets() never leads to buffer overflow.
- In the example:
 - If the line contains at most 998 characters, the entire line and the new-line character will be read and stored in the string.
 - If the line contains 999 or more characters, only the first 999 characters will be read and stored in the string.

Closing a file

- We can close a file simply by calling `fclose()` with the file pointer

```
FILE *fp;  
fp = (FILE *)fopen("new.txt", "w");  
if(fp == NULL){  
    printf("Error in file creation\n");  
    exit(0);  
}  
fprintf(fp, "Write whatever you want!\n");  
...  
fclose(fp);
```

Three special file streams

- Three special file streams are defined in the header file <stdio.h>. These FILE pointers are automatically opened in every program
 - `stdin` reads input from the keyboard
 - `stdout` send output to the screen
 - `stderr` prints errors to an error device (usually also the screen)
- `scanf(...)` is the same as `fscanf(stdin, ...)`
- `printf(...)` is the same as `fprintf(stdout, ...)`

Example

```
#include<stdio.h>
int main(){
    int x;

    fprintf(stdout,"Enter x\n");
    fscanf(stdin,"%d",&x);
    fprintf(stdout,"x: %d\n",x);
    fprintf(stderr,"No-error\n");
    return 0;
}
```

Output:

Enter x

43

x: 43

No-error

Reading and writing a character

- A character reading/writing is equivalent to reading/writing a byte

- `int getchar();`
- `int putchar(int c);`
- `int fgetc(FILE *fp);`
- `int fputc(int c, FILE *);`

```
char c;  
c = getchar();  
putchar(c);
```

```
char c;  
FILE *fp;  
c = fgetc(fp);  
fputc(c, fp);
```

Random access using fseek()

- `fseek()` – returns the present position of the file pointer

```
long ftell(FILE *fp)
```

- `fseek()` – can be used to set the position of a file pointer

```
int fseek(FILE *fp, long offset, int from_where)
```

- The new position, measured in bytes, is obtained by adding offset bytes to the position specified by `from_where`
- `from_where` can take one the following values
 - `SEEK_END` – end of the file
 - `SEEK_SET` – beginning of the file
 - `SEEK_CUR` – current position of the file pointer

Example

```
int main(){
    char c; FILE *fp;
    fp = fopen("input.txt", "r+");
    printf("%ld \n", ftell(fp));
    c = fgetc(fp); c = fgetc(fp);
    printf("%ld \n", ftell(fp));
    fseek(fp, 2, SEEK_CUR);
    printf("%ld \n", ftell(fp));
    fputs("fast purple",fp);
    printf("%ld \n", ftell(fp));
    fclose(fp);
    return 0;
}
```

Output:

```
0
2
4
15
```

input.txt:

Before: the quick brown fox jumped over the lazy dogs
After: the fast purple fox jumped over the lazy dogs

Example: fseek()

- Assume that we have file having 1000 prime numbers, one number in each line and there are no extra spaces
- We open the file in the read mode.
- We do not read the file from beginning to end.
- When the user specifies some n in the range [1,1000], we go to the location where the n -th prime is stored, and read that prime.
- 2 is the first prime (not the zero-th prime).
- We need to know exactly where the n -th prime is stored.
- The following statistics help us do that.
 - There are exactly four 1-digit primes.
 - There are exactly 21 2-digit primes.
 - There are exactly 143 3-digit primes.
 - The 1000-th prime is a 4-digit prime.
- We must not forget the new line character at the end of each line.

primes.txt:

2
3
5
7
11
...
107
...
7919

Example: fseek()

```
#include <stdio.h>
#include <stdlib.h>
int main() {
    FILE *fp; int n, p;
    fp = (FILE *)fopen("primes.txt", "r");
    while (1) {
        printf("Which prime? ");
        scanf("%d", &n);
        if((n < 1) || (n > 1000)) break;
        if(n <= 4) fseek(fp, (n - 1) * 2, SEEK_SET);
        else if(n <= 25) fseek(fp, 4 * 2 + (n - 5) * 3, SEEK_SET);
        else if(n <= 168) fseek(fp, 4 * 2 + 21 * 3 + (n - 26) * 4, SEEK_SET);
        else fseek(fp, 4 * 2 + 21 * 3 + 143 * 4 + (n - 169) * 5, SEEK_SET);
        fscanf(fp, "%d", &p); printf("%d-th prime is %d\n", n, p);
    }
    fclose(fp); exit(0);
}
```

Example: Copying a file

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
int main() {
    FILE *ifp, *ofp; int i, c; char srcfile[100], dstfile[100];
    strcpy(srcfile, "source.txt"); strcpy(dstfile, "copy.txt");
    if((ifp = (FILE *)fopen(srcfile, "r"))== NULL){
        printf("File does not exist\n"); exit(0);
    }
    if((ofp = (FILE *)fopen(dstfile, "w"))== NULL){
        printf("Unable to create file\n"); exit(0);
    }
    while ( (c = fgetc(ifp)) != EOF) fputc(c, ofp);
    fclose(ifp); fclose(ofp);
    return 0;
}
```

Example: reading & writing binary files

- We want to store Fibonacci numbers $F(0), F(1), \dots, F(40)$ in a file
- Text mode
 - We store 0, 1, 1, 2, 3, . . . , 102334155, one in a single line.
 - Some separators are needed between consecutive Fibonacci numbers (here we use '\n').
 - `fseek()` to locate $F(i)$ will be difficult.
- Binary mode
 - Assume that `int` is of size 32 bits (4 bytes).
 - We store the raw 32 bits of each $F(i)$ one after another.
 - No separators are needed because each $F(i)$ occupies exactly 4 bytes.
 - `fseek()` to locate $F(i)$ will be easy: just go to the $4i$ -th byte from the beginning.

Storing in human-readable format

```
#include <stdio.h>
#include <stdlib.h>
int main() {
    FILE *fp;
    int n = 40, i, F[41];
    F[0] = 0; F[1] = 1;
    for(i=2; i<=n; ++i) F[i] = F[i-1] + F[i-2];
    fp = (FILE *)fopen("Fib.txt", "w");
    if(fp == NULL) exit(1);
    for(i=0; i<=n; ++i) fprintf(fp, "%d\n", F[i]);
    fclose(fp);
    exit(0);
}
```

Storing in binary format

```
#include <stdio.h>
#include <stdlib.h>
int main() {
    FILE *fp; char *p;
    int n = 40, i, j, F[41];
    F[0] = 0; F[1] = 1;
    for(i=2; i<=n; ++i) F[i] = F[i-1] + F[i-2];
    fp = (FILE *)fopen("FibB.dat", "w");
    if(fp == NULL) exit(1);
    for(i=0; i<=n; ++i) {
        p = (char *)(F + i);
        for(j=0; j<4; ++j) { fprintf(fp, "%c", *p); ++p; }
    }
    fclose(fp);
    exit(0);
}
```

Reading in binary format

```
#include <stdio.h>
#include <stdlib.h>
int main() {
    FILE *fp; char *p;
    int n = 40, i, j, F[41];
    fp = (FILE *)fopen("FibB.dat", "r");
    for(i=0; i<=n; ++i) {
        p = (char *)(F+i);
        for(j=0; j<4; ++j) { fscanf(fp, "%c", p); ++p; }
    }
    fclose(fp);
    for(i=0; i<=n; ++i) printf("F(%d) = %d\n", i, F[i]);
    exit(0);
}
```

FibH.dat: 220 bytes
FibB.dat: 164 bytes

Redirection: Input & Output file

- One may redirect the standard input and output to other files, other than stdin or stdout
- Suppose, we have an executable code a.out

```
$ ./a.out < in.txt > out.txt
```

- scanf() will read data from the file “in.txt” and
- printf() will print the results on the file “out.txt”

- Another option

```
$ ./a.out < in.txt >> out.txt
```

- scanf() will read data from the file “in.txt” and
- printf() will append the results at the end of file “out.txt”

Command line arguments

- A program can be executed by directly typing the command at the shell prompt

```
$> gcc filename.c
```

```
$> ./a.out in.txt out.txt
```

```
$> exe_name param1 param2 ...
```

- The individual items specified are separated from one another by spaces. Use quotes to enter arguments with spaces.

- First item is the executable name

- Variables `argc` and `argv` keep track of the items specified in the command line

- Command line arguments may be passed by specifying them under `main()`

```
int main(int argc, char *argv[]);
```

argc - argument count, argv - array of string as command line arguments,
argv [] is NULL-terminated

Example

- \$> ./a.out in.txt out.txt
 - argc = 3
 - argv[0]="./a.out", argv[1]="in.txt", argv[2]="out.txt", argv[3]=NULL

Copying file: ./a.out <src> <dst>

```
int main(int argc, char *argv[]) {  
    FILE *ifp, *ofp; int i, c; char src[100], dst[100];  
    if(argc!=3){ printf("Usage: ./a.out <src> <dst>\n"); exit(0); }  
    else{ strcpy(src, argv[1]); strcpy(dst, argv[2]); }  
    if((ifp = (FILE *)fopen(src, "r"))== NULL){  
        printf("File does not exist\n"); exit(0);  
    }  
    if((ofp = (FILE *)fopen(dst, "w"))== NULL){  
        printf("Unable to create file\n"); exit(0);  
    }  
    while ( (c = fgetc(ifp)) != EOF) fputc(c, ofp);  
    fclose(ifp); fclose(ofp);  
    return 0;  
}
```

Useful library functions

- \$> ./a.out xyz 123 4.56 "fname lname"
- `strcpy(str, argv[1]);` — copies string
- `val = atoi(argv[2]);` — converts string to integer
- `cpi = atof(argv[3]);` — converts string to double
- `strcpy(name, argv[4]);` — copies string

Practice problems

- Write a program that reads a file, converts all lower-case letters to the upper case, and keeps the other characters intact, and stores the output in another file.
- Write a program that reads a 2-d array of integers from a file and replaces the contents of the file with the transpose of the matrix represented by the 2-d array.
- Write a program that reads student records containing name (string), roll_number (int), CGPA (float) from the user and writes them in a file, one record per line.
- Write a program that reads the records written by the above program into an array of structures. The structure should contain name, roll_number and CGPA as members.